

# ORM PERFORMANCE

---



# Performance Considerations

- Object Relational Mapping tools provide an increase in productivity
- Removes the necessity of “handcrafting” boilerplate\*\* SQL code.

## **HOWEVER**

- Lack of “insight” into the interactions with the database can reduce performance
- Without proper consideration of performance, scalability is also restricted.

## **On the other hand**

- A good ORM Framework has resources available for optimization.
- With the right approach, you can improve performance **BUT**

## **Ultimately You Need a DBA Expert**

Someone on the team who understands how a RDB works [Database design].

**Knowing SQL – is not database design**

\*\* rolled steel for making boilers



# Approach to Performance

**You should not do preemptive optimization**

Within reason, you should not care about optimization unless you **need** to. You need to **after** you have implemented, done measurements and found issues.

You should identify the scenarios that cause problems and **then** optimize to fix them.

**Obviously the focus of improvement is fetching relationships.**

**Obviously**

**The default Fetch strategy is `fetchType.LAZY`**

**Although**

Better performance comes from eager loading relationships that are always used

**REMEMBER:**

Implementation involves functionality & code stability

**Performance tuning comes AFTER**



# Hibernate FETCH Strategy

## *Hibernate NOT JPA*

- **Select fetching:** a second SELECT [ per parent N] is used to retrieve the associated collection. [**DEFAULT**] [N+1 Fetches] @Fetch(FetchMode.**SELECT**)
- **Join fetching:** associated collections are retrieved in the same SELECT, using an OUTER JOIN. [ 1 Fetch] [**EAGER**] @Fetch(FetchMode.**JOIN**)
- **Subselect fetching:** a second SELECT is used to retrieve the associated collections for all entities retrieved in a previous query or fetch. [2 Fetches]
- @Fetch(FetchMode.**SUBSELECT**)
- **Batch fetching:** Optimization of Select Fetching. Associated collections are fetched according to declared Batch Size(N/Batch Size) + 1; @BatchSize(size=n)
- **EXAMPLE:**
- @Fetch(FetchMode.**SUBSELECT**)
- **private** Set<Address> addresses;



# Hibernate Fetch Strategy Issues

## Select

- N+1 TBA [To Be Avoided]

## Join

Cartesian – need to watch collection sizes; can be useful strategy

## SubSelect

depends on the “parent” query. If parent Query is complex, it could have performance impacts.

If fetch=FetchType.**LAZY** need to “hydrate” children

## BatchSize

# of Fetches “unknown” UNLESS size of parent is constant

Batch fetching is often called a blind-guess optimization

If fetch=FetchType.**LAZY** need to “hydrate” children

REMEMBER: fetch=FetchType.**LAZY** is recommended Default




## Example N = 3

# N+1 Problem

- Member has a OneToMany relationship with Address

Uses "Default" strategy of SELECT



- Declared as Fetch EAGER:

`@OneToMany(mappedBy="member", fetch=FetchType.EAGER)`

`private Set<Address> addresses = new HashSet<Address>();`

- For

`entityManager.createQuery( "from Member"`

`.getResultList();`

ORM will issue ONE fetch for ALL the Members

&

And N fetches; one for each member's Address Collection

- So for 3 members we have ONE fetch for Members and THREE fetches for Address collections

### N+1 ISSUE

```

Hibernate:
select
  member0_.member_id as member_i1_5_,
  member0_.age as age2_5_,
  member0_.firstName as firstNam3_5_,
  member0_.lastName as lastName4_5_,
  member0_.memberNumber as memberNu5_5_,
  member0_.title as title6_5_
from
  Member member0_
Hibernate:
select
  addresses0_.member_id as member_i6_5_0_,
  addresses0_.id as id1_0_0_,
  addresses0_.id as id1_0_1_,
  addresses0_.city as city2_0_1_,
  addresses0_.member_id as member_i6_0_1_,
  addresses0_.state as state3_0_1_,
  addresses0_.street as street4_0_1_,
  addresses0_.zipCode as zipCode5_0_1_
from
  Address addresses0_
where
  addresses0_.member_id=?
Hibernate:
select
  addresses0_.member_id as member_i6_5_0_,
  addresses0_.id as id1_0_0_,
  addresses0_.id as id1_0_1_,
  addresses0_.city as city2_0_1_,
  addresses0_.member_id as member_i6_0_1_,
  addresses0_.state as state3_0_1_,
  addresses0_.street as street4_0_1_,
  addresses0_.zipCode as zipCode5_0_1_
from
  Address addresses0_
where
  addresses0_.member_id=?
Hibernate:
select
  addresses0_.member_id as member_i6_5_0_,
  addresses0_.id as id1_0_0_,
  addresses0_.id as id1_0_1_,
  addresses0_.city as city2_0_1_,
  addresses0_.member_id as member_i6_0_1_,
  addresses0_.state as state3_0_1_,
  addresses0_.street as street4_0_1_,
  addresses0_.zipCode as zipCode5_0_1_
from
  Address addresses0_
where
  addresses0_.member_id=?
  
```



# Join Fetching Strategy

Hibernate Annotation: `@Fetch(FetchMode.JOIN)`

**ALWAYS**

Causes an **EAGER** fetch of the child collections

This is because the characteristic of a Join is ONE fetch  
Parent & Child TOGETHER

This can only be accomplished by loading the child  
collection when the parent is fetched [ ~= **EAGER** fetch]

**Best practice: Use Lazy Initialization [NOT Eager]**

So we will implement the Join Fetch “Manually”



# Join == Cartesian “LIKE” Product Problem

- For sets A and B, the Cartesian product is  $A \times B$
- For sets A,B and C, the Cartesian product is  $A \times B \times C$  - etc.

Member has a OneToMany relationship with Address

```
@OneToMany(mappedBy="member", fetch=FetchType.LAZY)
```

```
private Set<Address> addresses = new HashSet<Address>();
```

For:

```
Query query=entityManager.createQuery("SELECT m
FROM Member AS m JOIN FETCH m.addresses AS a");
```

ORM will do **ONE FETCH** BUT will generate duplicates

]

ORM NOTE: “Product” =

$\sum$  # Addresses [per Member.

Sean has 2 Addresses so 2 copies ; Bill has 3 so 3 copies

N+1 GONE - Join Fetch - Cartesian Product  
Hibernate:

```
select
  member0.member_id as member_i1_5_0_,
  addresses1.id as id1_0_1_,
  member0.age as age2_5_0_,
  member0.firstName as firstNam3_5_0_,
  member0.lastName as lastName4_5_0_,
  member0.memberNumber as memberNu5_5_0_,
  member0.title as title6_5_0_,
  addresses1.city as city2_0_1_,
  addresses1.member_id as member_i6_0_1_,
  addresses1.state as state3_0_1_,
  addresses1.street as street4_0_1_,
  addresses1.zipCode as zipCode5_0_1_,
  addresses1.member_id as member_i6_5_0_1_,
  addresses1.id as id1_0_0_1_
from
  Member member0_
inner join
  Address addresses1_
on member0_.member_id=addresses1.member_id
```

Member Name :	Sean	Smith
Address :	Batavia	Iowa
Address :	Red Rock	Iowa
Member Name :	Sean	Smith
Address :	Batavia	Iowa
Address :	Red Rock	Iowa
Member Name :	Bill	Due
Address :	Washington	Iowa
Address :	Mexico	Iowa
Address :	Paris	Iowa
Member Name :	Bill	Due
Address :	Washington	Iowa
Address :	Mexico	Iowa
Address :	Paris	Iowa
Member Name :	Bill	Due
Address :	Washington	Iowa
Address :	Mexico	Iowa
Address :	Paris	Iowa



# Join == Cartesian Product Problem

For sets A and B, the Cartesian product is  $A \times B$

For sets A,B and C, the Cartesian product is  $A \times B \times C$  - etc.

Member has a OneToMany relationship with Address & Order

```
@OneToMany(mappedBy="member", fetch=FetchType.LAZY)
```

```
private Set<Address> addresses = new HashSet<Address>();
```

```
@OneToMany(mappedBy="member", fetch=FetchType.LAZY)
```

```
private Set<Order> orders = new HashSet<Order>();
```

For:

```
Query query = entityManager.createQuery("SELECT m FROM Member
    AS m JOIN FETCH m.orders as o
    JOIN FETCH m.addresses as a where m.id = :id");
```

ORM will do **ONE FETCH** BUT will generate duplicates

**ORM NOTE: Product = 2 Addresses X 3 Order  
SIX Members**

```
Member Name : Sean Smith
Address : Batavia Iowa
Address : Red Rock Iowa
Order : Order 3
Order : Order 2
Order : Order 1

Member Name : Sean Smith
Address : Batavia Iowa
Address : Red Rock Iowa
Order : Order 3
Order : Order 2
Order : Order 1

Member Name : Sean Smith
Address : Batavia Iowa
Address : Red Rock Iowa
Order : Order 3
Order : Order 2
Order : Order 1

Member Name : Sean Smith
Address : Batavia Iowa
Address : Red Rock Iowa
Order : Order 3
Order : Order 2
Order : Order 1

Member Name : Sean Smith
Address : Batavia Iowa
Address : Red Rock Iowa
Order : Order 3
Order : Order 2
Order : Order 1
```



# Clean Up The Cartesian Product Data

- Query `query=entityManager.createQuery("SELECT DISTINCT m  
FROM Member AS m JOIN FETCH m.addresses AS a");`

- DISTINCT keyword removes duplicates

**However**

**It accomplishes it in Memory [ After DB fetch]**

```
Hibernate:
select
    distinct member0_.member_id as member_id_0_1_,
    addresses1_.id as id1_0_1_,
    member0_.age as age2_5_0_,
    member0_.firstName as firstName3_5_0_,
    member0_.lastName as lastName4_5_0_,
    member0_.memberNumber as memberNumber5_5_0_,
    member0_.title as title6_5_0_,
    addresses1_.city as city2_0_1_,
    addresses1_.member_id as member_id_6_0_1_,
    addresses1_.state as state3_0_1_,
    addresses1_.street as street4_0_1_,
    addresses1_.zipCode as zipCode5_0_1_,
    addresses1_.member_id as member_id_6_5_0_1_,
    addresses1_.id as id1_0_0_1_
from
    Member member0_
inner join
    Address addresses1_
        on member0_.member_id=addresses1_.member_id
Member Name.: Sean Smith
Address.: Batavia Iowa
Address.: Red Rock Iowa
Member Name.: Bill Due
Address.: Washington Iowa
Address.: Mexico Iowa
Address.: Paris Iowa
```



# SubSelect Fetch

- Declared as Fetch LAZY:
- `@OneToMany(mappedBy="member", fetch=FetchType.LAZY)`
- `@Fetch(FetchMode.SUBSELECT)`
- `private Set<Address> addresses = new HashSet<Address>();`

For `FetchMode.SUBSELECT`

ORM will do ONE Fetch for All Parents

ORM will do ONE Fetch for All child collections

- Need to “Hydrate” collections:

```
List<Member> members = (List<Member>)this.findAll();
members.get(0).getAddresses().get(0);
```

We can also do this “Manually” WITHOUT

`@Fetch(FetchMode.SUBSELECT)`

*See DEMO FetchSubSelect*

Hibernate:  
select  
member0\_.member\_id as m  
member0\_.age as age2\_5  
member0\_.firstName as f  
member0\_.lastName as la  
member0\_.memberNumber a  
member0\_.title as title  
from  
Member member0\_

Hibernate:  
select  
addresses0\_.member\_id a  
addresses0\_.id as id1\_0  
addresses0\_.id as id1\_0  
addresses0\_.city as cit  
addresses0\_.member\_id a  
addresses0\_.state as st  
addresses0\_.street as s  
addresses0\_.zipCode as  
from  
Address addresses0\_  
where  
addresses0\_.member\_id i  
?, ?, ?  
)

Batch fetch example  
Member Name : Sean Smith  
Address : Red Rock Iowa  
Address : Batavia Iowa  
Member Name : Peat Moss  
Member Name : Bill Due  
Address : Mexico Iowa  
Address : Washington Iowa  
Address : Paris Iowa



# Batch Size Fetch

- Declared as Fetch LAZY:
- `@OneToMany(mappedBy="member", fetch=FetchType.LAZY)`
- `@Fetch(FetchMode.SELECT)`
- `@BatchSize(size = 3)`
- `private Set<Address> addresses = new HashSet<Address>();`
- Need to "Hydrate" collections:
- `List<Member> members = (List<Member>)this.findAll();`
- `for (Member member : members)`
- `if (!member.getAddresses().isEmpty())`
- `member.getAddresses().get(0);`
- In example, ORM will do ONE Collection Fetch  
[based on batch size = # parents]

## NOTE:

In example # members = 3; BatchSize = 3; 1 Collection fetch  
 If # members = 3; BatchSize = 2; 2 Collection fetches  
 If # members = 3; BatchSize = 1; 3 Collection fetches

Hibernate:

```
select
  member0_.member_id as m
  member0_.age as age2_5_
  member0_.firstName as f
  member0_.lastName as la
  member0_.memberNumber a
  member0_.title as title
from
  Member member0_
```

Hibernate:

```
select
  addresses0_.member_id a
  addresses0_.id as id1_0_
  addresses0_.id as id1_0_
  addresses0_.city as cit
  addresses0_.member_id a
  addresses0_.state as st
  addresses0_.street as s
  addresses0_.zipCode as
from
  Address addresses0_
where
  addresses0_.member_id i
  ?, ?, ?
)
```

Batch fetch example

```
Member Name : Sean Smith
Address : Red Rock Iowa
Address : Batavia Iowa
Member Name : Peat Moss
Member Name : Bill Due
Address : Mexico Iowa
Address : Washington Iowa
Address : Paris Iowa
```



# Customized Fetching Strategies

**HOW** to Access an Association [for Performance Tuning]

**WHEN** is **FetchType.LAZY** OR FetchType.**EAGER**

## The Default Fetching Strategy

is not used to customize fetching [ it is not efficient]

## Use & Master The Query API

to override the Default Fetching Strategy for custom fetching

## Understand & Use The Join fetch & SubSelect

Avoid the N+1 fetch problem as a rule of thumb

FetchMode.**SELECT**(Default) is extremely vulnerable to the N+1 problem

## Use Lazy Initialization

- **Regardless of your fetching strategy WHEN you fetch takes “priority”**
- \*\*\* Non-lazy initialization [FetchType.**EAGER** ] **always** loads the relationship \*\*\*

**REMEMBER - Optimize Selectively**



# Main Point

- An eager fetch strategy often creates inefficient queries. To increase performance we can configure the entity associations with a lazy fetch strategy and fetch them in a more efficient manner.
- ***Science of Consciousness:*** *As we grow in creative intelligence we enjoy the advantage of increased efficient action..*



# Second Level Cache

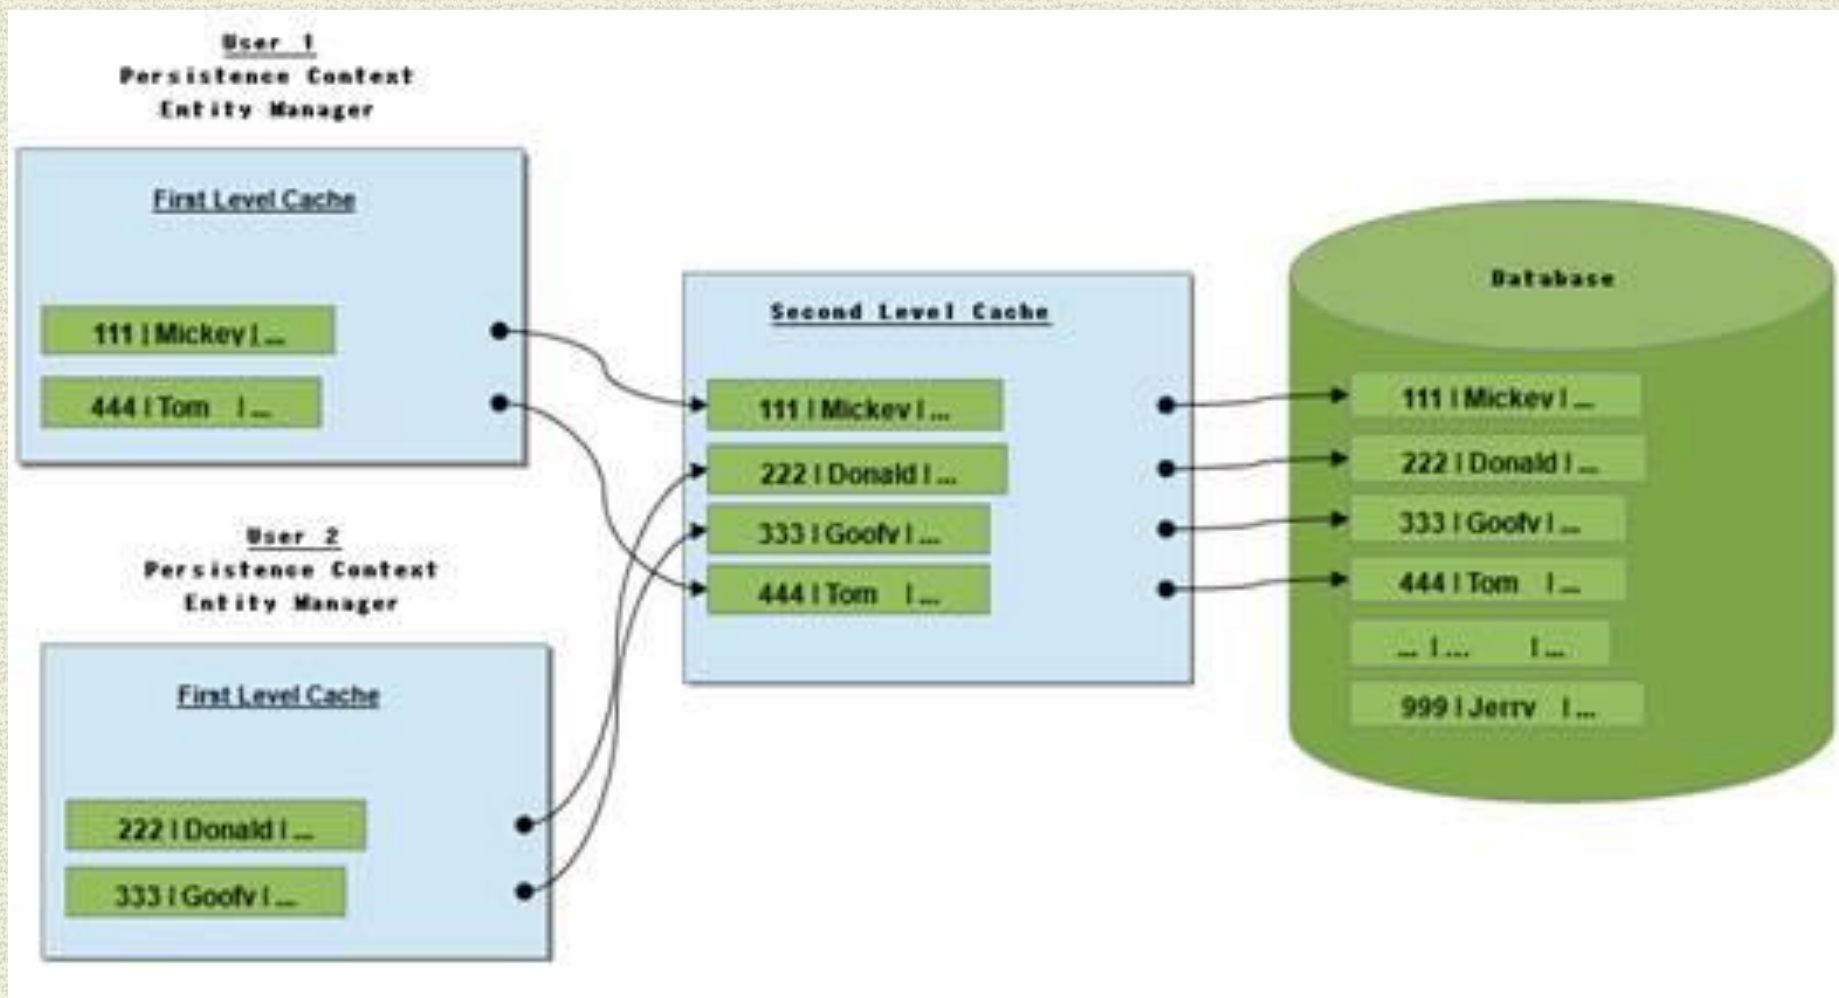
- **A second-level cache**

A local store of entity data managed by the persistence provider to improve application performance

- Improves performance by avoiding expensive database calls
- Keeps the entity data local to the application
- Transparent to the application
- Available across all users [ Application wide]
- Complements First level cache



# Level 2 Cache





# Query for Entity

- Check First level Cache
- If Found:
  - Return Entity
- If not Found:
  - Check Second level Cache
  - If Found:
    - Update First Level Cache
    - Return Entity
  - If not Found:
    - Execute DB Query
    - Update Second Level Cache
    - Update First Level Cache
    - Return Entity



# Cache Concurrency Strategy

- **Transactional:**

**Read-Mostly** data where it is critical to prevent stale data in concurrent transactions, in the rare case of an update.

***Full Lock on all entities in transaction. Performance issue.***

- **Read-write:**

**Read-Mostly** data where it is critical to prevent stale data in concurrent transactions, in the rare case of an update.

Supports Repeatable Read Isolation - Allows phantom reads

- **Nonstrict-read-write:**

No guarantee of consistency between the cache and the database.

Use this strategy if data hardly ever changes and a small likelihood of stale data is not of critical concern.

- **Read-only:**

Suitable for data which never changes. Use it for reference data only.

***Simplest and optimal performing strategy.***



# READ Only – Low Hanging Fruit



- **Read-only caches are easy to handle**

It is immutable (Modification Forbidden )

No consistency issues.

**Always** a good candidate for second level caching

- **Read-write caches are more “subtle” in their behavior.**

- Interaction with the Hibernate session can lead to unwanted behavior.

- The benefits of the C in ACID are compromised if cache is out of sync with DB

**Eventual consistency is NOT a primary use case of Relational DBs in an Enterprise.**



# Second Level Cache Decision

**Database queries slow?**

**Second Level Cache is the *final resort***

- **Optimize ORM Queries**

Make sure the fetching strategy is properly designed,

Remove N+1 query problems

Involve the DBA [Expert]

Employ indexes

Investigate data base solutions [ e.g. partitioning]

If performance is still an issue **THEN** consider a second level cache



# Two Types of Cache

## 1. **Second Level Cache**

Second level cache is a key-value store.

Applicable for accessing an entity by Primary Key

### ***These will hit Second Level Cache***

```
member = memberService.findOne(1L);  
entityManager.createQuery("SELECT m FROM Member m where  
    m.id= :member");
```

### ***These will NOT hit cache***

```
entityManager.createQuery("SELECT m FROM Member AS m JOIN FETCH  
    m.addresses AS a where m.id= :member");  
entityManager.createQuery("select m from Member m where  
    m.memberNumber =:number");
```

## 2. **Query Cache**

Applicable for accessing entities by specific query

**The two preceding queries are Query Cache “candidates”**

**NOTE: Query cache store the query in query cache and the actually entities in the Second Level Cache**



# Configuring resources for Caching

For Second Level Cache:

```
@Cache(usage = org.hibernate.annotations.CacheConcurrencyStrategy.READ_ONLY)  
public class Member {
```

For Query Cache:

```
@NamedQuery(name="Member.findById",  
            query="select m from Member m where m.memberNumber = :memberNumber ",  
            hints={ @QueryHint(name="org.hibernate.cacheable", value="true") } )
```



# Configuration

- **pom.xml**
- `<dependency>`
- `<groupId>net.sf.ehcache</groupId>`
- `<artifactId>ehcache</artifactId>`
- `</dependency>`
- `<dependency>`
- `<groupId>org.hibernate</groupId>`
- `<artifactId>hibernate-ehcache</artifactId>`
- `</dependency>`
- **Hibernate Properties**
- `<prop key="hibernate.cache.provider_class">`  
         `org.hibernate.cache.EhCacheProvider</prop>`
- `<prop key="hibernate.cache.region.factory_class">`  
         `org.hibernate.cache.ehcache.SingletonEhCacheRegionFactory</prop>`
- `<prop key="hibernate.cache.use_second_level_cache">true</prop>`
- `<prop key="hibernate.cache.use_query_cache">true</prop>`
-



# ehCache.xml

<ehcache>

```
<diskStore path="java.io.tmpdir"/>
```

```
<defaultCache
```

```
    maxElementsInMemory="1000"  
    eternal="false"  
    timeToIdleSeconds="120"  
    timeToLiveSeconds="120"  
    overflowToDisk="true"  
/>
```

```
<cache name="edu.mum.domain.Member"
```

```
    maxElementsInMemory="1000"  
    eternal="false"  
    timeToIdleSeconds="600"  
    timeToLiveSeconds="3600"  
    overflowToDisk="true"  
/>
```

```
<cache name="edu.mum.domain.Address"
```

```
    maxElementsInMemory="1000"  
    eternal="false"  
    timeToIdleSeconds="600"  
    timeToLiveSeconds="3600"  
    overflowToDisk="true"  
/>
```

</ehcache>

The defaultCache is applied to any cache not explicitly configured



# Spring Cache Abstraction

- Applies caching to **Java methods**

When method is invoked, the abstraction will apply a caching behavior checks whether the method has been ***already executed for the given arguments***.

[ Does **NOT** re-execute the method a second time...]

- **Will Work across “Query” methods**

Does not distinguish between “Second Level Cache” & Query Cache  
Handles Both

- Spring Cache Abstraction is separate from Hibernate cache
- Hibernate Second Level Cache and Spring Cache (method caching) can co-exist with Ehcache as the underlying provider for both.



# Cache Abstraction Annotations


- **The abstraction provides following Java annotations:**
- **@Cacheable:**  
Put the method returned value(s) into the cache [ **READ ONLY**]
- **@CacheEvict:**  
Remove an entry from the cache [stale or unused...]
- **@CachePut:**  
Force a cache entry to be updated [ **READ WRITE related**]



# Spring Cache Configuration

- **Declaration** – refers to ehCache.xml resource

- `@Cacheable("edu.mum.domain.Member")`
- `public Member findOne(Long id) {`
- `return memberDao.findOne(id);`
- `}`



```
<cache name="edu.mum.domain.Member"
      maxElementsInMemory="1000"
      eternal="false"
      timeToIdleSeconds="600"
      timeToLiveSeconds="3600"
      overflowToDisk="true"
      />
```

- **XML Configuration**

- `<cache:annotation-driven />`
- `<bean id="cacheManager"`  
`class="org.springframework.cache.ehcache.EhCacheCacheManager">`
- `<constructor-arg ref="ehcacheManager" />`
- `</bean>`
- `<bean id="ehcacheManager"`  
`class="org.springframework.cache.ehcache.EhCacheManagerFactoryBean">`
- `<property name="configLocation" value="ehcache.xml" />`
- `</bean>`



# Main Point

1. The second level cache can permit more scalability for the application, when the same entities are retrieved by multiple users they could be retained in memory to avoid subsequent accesses to the database, increasing efficiency & speed.
2. ***Science of Consciousness:*** *Through the constant practice of Transcendental Meditation, the reactions of mind and body are faster.*



